

Attention with a worked out example

We want the word **walk** to be more closely associated to **birds** in this specific text:

Birds sometimes walk

We want to do this because there are other examples in text:

Sometimes for walking birds use floppy feet...

OR But birds don't just walk they fly...

Lets say we have a poorly trained Value vector for **V**:

$$V = \begin{bmatrix} 0.1 & 0.2 & 0.3 \\ -0.8 & 0.3 & -0.1 \\ 0.5 & -0.1 & -0.4 \end{bmatrix}$$

where each token has a representation like:

$$\begin{aligned} v1(Birds) &\rightarrow [0.1, 0.2, 0.3] \\ v2(sometimes) &\rightarrow [-0.8, 0.3, -0.1] \\ v3(Walk) &\rightarrow [0.5, -0.1, -0.4] \end{aligned}$$

and say you have a reasonably good attention **A** matrix trained

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 0.1 & 0.9 & 0 \\ 0.6 & 0.1 & 0.4 \end{bmatrix}$$

Initially the dot product of the representations is pretty low:

$$\vec{v_3} \cdot \vec{v_1} = 0.1 * 0.5 + 0.2 * -0.1 + 0.3 * -0.4 = 0.05 - 0.02 - 0.12 = \mathbf{-0.09}$$

Now we want **walk** to pay more attention to **bird** and **sometimes** from a learnt **A**

$$\hat{v}_3 = \begin{bmatrix} a_{3,1} * v_{11} + a_{3,2} * v_{21} + a_{3,3} * v_{31} \\ a_{3,1} * v_{12} + a_{3,2} * v_{22} + a_{3,3} * v_{32} \\ a_{3,1} * v_{13} + a_{3,2} * v_{23} + a_{3,3} * v_{33} \end{bmatrix}$$

$$\hat{v}_3 = \begin{bmatrix} a_{31} * v_1 + a_{32} * v_2 + a_{33} * v_3 & \text{(Attention paid by 'walk' to 'bird' and 'sometimes' *across* first column)} \\ a_{31} * v_1 + a_{32} * v_2 + a_{33} * v_3 & \text{(Same thing but across the second dimension)} \\ a_{31} * v_1 + a_{32} * v_2 + a_{33} * v_3 & \text{(... And the third dimension)} \end{bmatrix}$$

$$\hat{v}_3 = \begin{bmatrix} 0.6 * 0.1 + 0.1 * -0.8 + 0.4 * 0.5 \\ 0.6 * 0.2 + 0.1 * 0.3 + 0.4 * -0.1 \\ 0.6 * 0.3 + 0.1 * -0.1 + 0.4 * -0.4 \end{bmatrix}$$

$$\hat{v}_3 = \begin{bmatrix} 0.06 - 0.08 + 0.2 \\ 0.12 + 0.03 - 0.04 \\ 0.18 - 0.01 - 0.16 \end{bmatrix}$$

It would seem that if we have learnt **A** well, i.e. **bird** is paying more attention to **walk**, then it seems to be now closer to **walk** in the embedded space.

$$\hat{v}_3 = \begin{bmatrix} 0.18 \\ 0.11 \\ 0.01 \end{bmatrix}$$

Now:

$$\hat{v}_3 \cdot v_3 = 0.18 * 0.1 + 0.11 * 0.2 + 0.01 * 0.3 = 0.018 + 0.022 + 0.003 = \mathbf{0.043}$$

i.e. a better association between **walk** and **bird** in the embedded space.

Thoughts on the Mathematics of this

1. I *think*, we are saying that the word **walk** is a **weighted** combination of all the other word embeddings. And therefore I would conjecture that this may not be enough, linear functions are good but nonlinear are better so **V** will also be forward propagated to an MLP to learn any non linearity that's left.
2. Also notice that the final vector are three terms and over **n-grams** of differing sequence length. So perhaps the vector v_3 terms are these piece-wise, auto regressive collection of terms each is a linear function. This kind of reminds me of interaction terms in linear regression and how indicator variables are used to capture non-linear relationships. Is this thought too clever by half in how I think about these things?

Thoughts on how does **bird** get attended to by **flying** in the same text where its associated with **walking**?

We might have other examples in the text of **bird** being associated with the words **fly** too in the text. How could the network learn about it and predict it? I think the answer is **capacity** which loosely translates to the number of parameters in the network we can use. Lets say we have a sentence:

But birds are good at ____.

And the word we likely want learnt is **flying**.

Capacity here would mean the opportunities for the word **flying** to learn the word association with the word **bird** and possibly **but** and **good**. Well if we consider just one attention head

1. There are **n_head** parameters for each vector in **V** opportunities to learn these associations. Is there a “partition” of dimension of learn’t **v_i** that is different for **walk** and **flying**?
2. Could it be that **walk** and **fly** i.e. the $a_{i,j}$ for i th token – which are learnt via **Q,K** – are different for their preceeding phrases "**birds sometimes ____**" and "**birds are good at ____**"?
3. In the Query matrix which has a **n_head** size vector Query for **flying**. Lets call this row **q_k**.
 - And Since we compute the attention row **a_i** as a product of **q_k** and **V** we have **n_head*n_head** parameters, which is typically 48,000 different parameters just for whatever **flying** needs to attend to! A big space indeed.
4. Not to mention more parameters in MLP and other heads that get learnt later.

Here is some matrix math for this stuff

$$A = \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & a_{i,j} & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} V = \begin{bmatrix} \cdot & v_0 & \cdot \\ \cdot & v_1 & \cdot \\ \cdot & v_2 & \cdot \end{bmatrix} \hat{V} = A.V$$

where V are the old embeddings and A are the Attention dot product.

$$\hat{v}_{i,j} = \sum_{k=1}^{d_k} a_{i,k} v_{k,j}$$

$$\hat{\underline{v}}_i = \sum_{j=1}^{d_k} a_{i,j} * \underline{v}_j$$

So we are *interacting* previously independent v_j s across the embedding.
